

Per-size sAP loss vs GPU contention (ResNet50 k-sweep)
All-NPU loss ≈ 0 (flat, not shown); rerun w/ util sampler

