





Steady-state runtime overhead of the safety layer



We run the same stable workload (`tests/steady\_state\_views\_schedule.yaml`) -- four view models on CPU at moderate rates

- **Throughput, per-view latency, and drop rate** are indistinguishable across the three modes once the run-to-run variance is
- **Process-tree CPU** is the most direct view of monitoring cost. Adaptive hot-swap adds +17.0 percentage points (+0.75% re
- **Take-away.** The safety layer (QoS monitoring + transition validation + rollback) imposes no measurable tax on view-level